

# Rainforest Hack Hour

*Andrew Lowe*

*21 March 2018*

## Business challenge

- Quantify the impact of various drivers on the share performance of a specific brand across a selected market.
- Predict the share performance of brands for future periods.

In the following, we will prove that we have achieved both objectives.

## Configure setup and load the data

The easiest way to import the data into R is to first export it from Spotfire as an Excel spreadsheet. (Unfortunately, exporting the data as a tab-delimited text file results in weird encoding issues that cause errors like “embedded nul in string” when using `read.table` or `fread` to read the data in R.):

```
suppressPackageStartupMessages( # Tidyverse is verbose on startup
  require(tidyverse) # We'll need tools from this library
)

## Warning: package 'tidyverse' was built under R version 3.4.2
## Warning: package 'tibble' was built under R version 3.4.4
## Warning: package 'tidyr' was built under R version 3.4.4
## Warning: package 'purrr' was built under R version 3.4.2
## Warning: package 'dplyr' was built under R version 3.4.2
## Warning: package 'stringr' was built under R version 3.4.3
## Warning: package 'forcats' was built under R version 3.4.4
set.seed(42) # Set random number seed for the sake of reproducibility
require(readxl) # For reading Excel files

## Loading required package: readxl
dat.xls <- read_xls("Data Table.xls") # Read in data
names(dat.xls) # Print column names

## [1] "Column 1" "TimeName"
## [3] "PeriodType" "PeriodEndDate"
## [5] "Region" "Country"
## [7] "Area" "AreaHierarchy"
## [9] "Category" "Company"
## [11] "Brand" "Form"
## [13] "Concentration" "SecondBenefit"
## [15] "NumberOfJobs" "BasicSize"
## [17] "Item" "ValueSalesMLC"
## [19] "VolumeSalesMSU" "UnitSalesx1000"
## [21] "WeightedDistribution" "WDDisplay"
```

```
## [23] "WDAnyPromo"           "WDFeature"
## [25] "WDPriceCut"           "NumSize"
## [27] "PPU"                  "PPSU"
## [29] "PricePerScoop"        "PromoPerDistribution (2)"
## [31] "ScoopsSold (2)"       "Fitted"
## [33] "Resid"                "Filtered out at 5:38:00 PM"
## [35] "ValidPpuCount"        "Date"
## [37] "PromoPerDistribution" "ScoopsSold"
## [39] "PriceCutPerDistribution"
```

```
View(dat.xls)
```

## Data cleaning

We subset the data to contain just the columns that we believe we need to build a predictive model:

```
# Select the columns we need (according to Tamas Sarkadi <Tamas_Sarkadi@epam.com>):
```

```
dat <- dat.xls[, names(dat.xls) %in% c(
  "Category",
  "Company",
  "Brand",
  "Form",
  "Concentration",
  "BasicSize",
  "SecondBenefit",
  "NumberOfJobs",
  "Item",
  "ValueSalesMLC", # Target
  "Date", # Timestamp
  "WeightedDistribution", # Feature
  "WDFeature", # Feature
  "WDDisplay", # Feature
  "WDPriceCut", # Feature
  "PPSU" # Feature
)]
names(dat)
```

```
## [1] "Category"           "Company"           "Brand"
## [4] "Form"               "Concentration"      "SecondBenefit"
## [7] "NumberOfJobs"       "BasicSize"         "Item"
## [10] "ValueSalesMLC"      "WeightedDistribution" "WDDisplay"
## [13] "WDFeature"          "WDPriceCut"        "PPSU"
## [16] "Date"
```

Some columns are single valued and therefore have no predictive value; they are removed:

```
strip.single.valued <- function(df) {
  only.one.value <- sapply(df, function(x) length(unique(x)) == 1) # Count unique values
  str(df[1, only.one.value]) # These columns only have one unique value! Print them
  df <- df[, !only.one.value] # Remove those columns
  return(df)
}

dat <- strip.single.valued(dat) # Do it
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame': 1 obs. of 3 variables:
```

```
## $ Category : chr "LAUNDRY DETERGENTS V2 CATEGORY"
## $ WDDisplay: logi NA
## $ WDFeature: logi NA
```

The date information (in *POSIXct* format) is transformed into *Date* format:

```
require(lubridate)
```

```
## Loading required package: lubridate
##
## Attaching package: 'lubridate'
## The following object is masked from 'package:base':
##
##     date
```

```
dat$Date <- ymd(dat$Date) # Transform into Date object
```

We remove any duplicated columns, if present:

```
# Remove columns with duplicate entries by fast comparison of hashes:
require(digest)
```

```
## Loading required package: digest
## Warning: package 'digest' was built under R version 3.4.3
duplicate.columns <- names(dat)[duplicated(lapply(dat, digest))]
if(length(duplicate.columns) == 0) {# Are there any duplicate columns?
  print("No duplicated columns")
} else {
  print(duplicate.columns)
}
```

```
## [1] "No duplicated columns"
```

```
dat <- dat[, !names(dat) %in% duplicate.columns]
names(dat)
```

```
## [1] "Company"          "Brand"            "Form"
## [4] "Concentration"    "SecondBenefit"    "NumberOfJobs"
## [7] "BasicSize"        "Item"             "ValueSalesMLC"
## [10] "WeightedDistribution" "WDPriceCut"       "PPSU"
## [13] "Date"
```

We transform character strings to categorical variables:

```
dat <- dat %>% mutate_if(is.character, as.factor)
```

Now we do the drill-down to get the products/SKUs. How many do we have?

```
dat %>% group_by(Company,
                 Brand,
                 Form,
                 Concentration,
                 SecondBenefit,
                 NumberOfJobs,
                 BasicSize,
                 Item) %>%
  mutate(SKU = paste( # Add SKU id
                    Company,
```

```

    Brand,
    Form,
    Concentration,
    SecondBenefit,
    NumberOfJobs,
    BasicSize,
    Item,
    sep = " | ")
  ) %>%
  arrange(SKU, Date) %>% # Order by SKU then Date
  mutate(count = n()) -> products

products %>% count() %>% nrow() # How many products?

## [1] 691

length(unique(products$SKU)) # Check maths: wow many SKUs? Should be identical.

```

```
## [1] 691
```

Some products have data reported for many time points, while others have little data:

```
products %>% pull(count) %>% summary() # Print summary of counts
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.0   78.0   113.0   125.7   181.0   258.0
```

What's the maximum size of the reporting period?

```
range(products$Date)
```

```
## [1] "2010-11-07" "2015-10-18"
```

Measurements appear to be taken weekly:

```
head(sort(unique(products$Date)))
```

```
## [1] "2010-11-07" "2010-11-14" "2010-11-21" "2010-11-28" "2010-12-05"
## [6] "2010-12-12"
```

We plot a quick glimpse of the data for the SKUs with more than the mean number of measurements.

```
products %>% filter(count > 125) %>%
  ggplot(aes(x = Date, y = ValueSalesMLC, group = SKU)) +
  facet_wrap(~SKU, scales = "free_y", ncol = 6) +
  geom_line() +
  theme_bw() +
  theme(
    strip.background = element_blank(),
    strip.text.x = element_blank(),
    axis.title.x = element_blank(),
    axis.title.y = element_blank(),
    axis.ticks.x = element_blank(),
    axis.ticks.y = element_blank(),
    axis.text.x = element_blank(),
    axis.text.y = element_blank()
  )

```

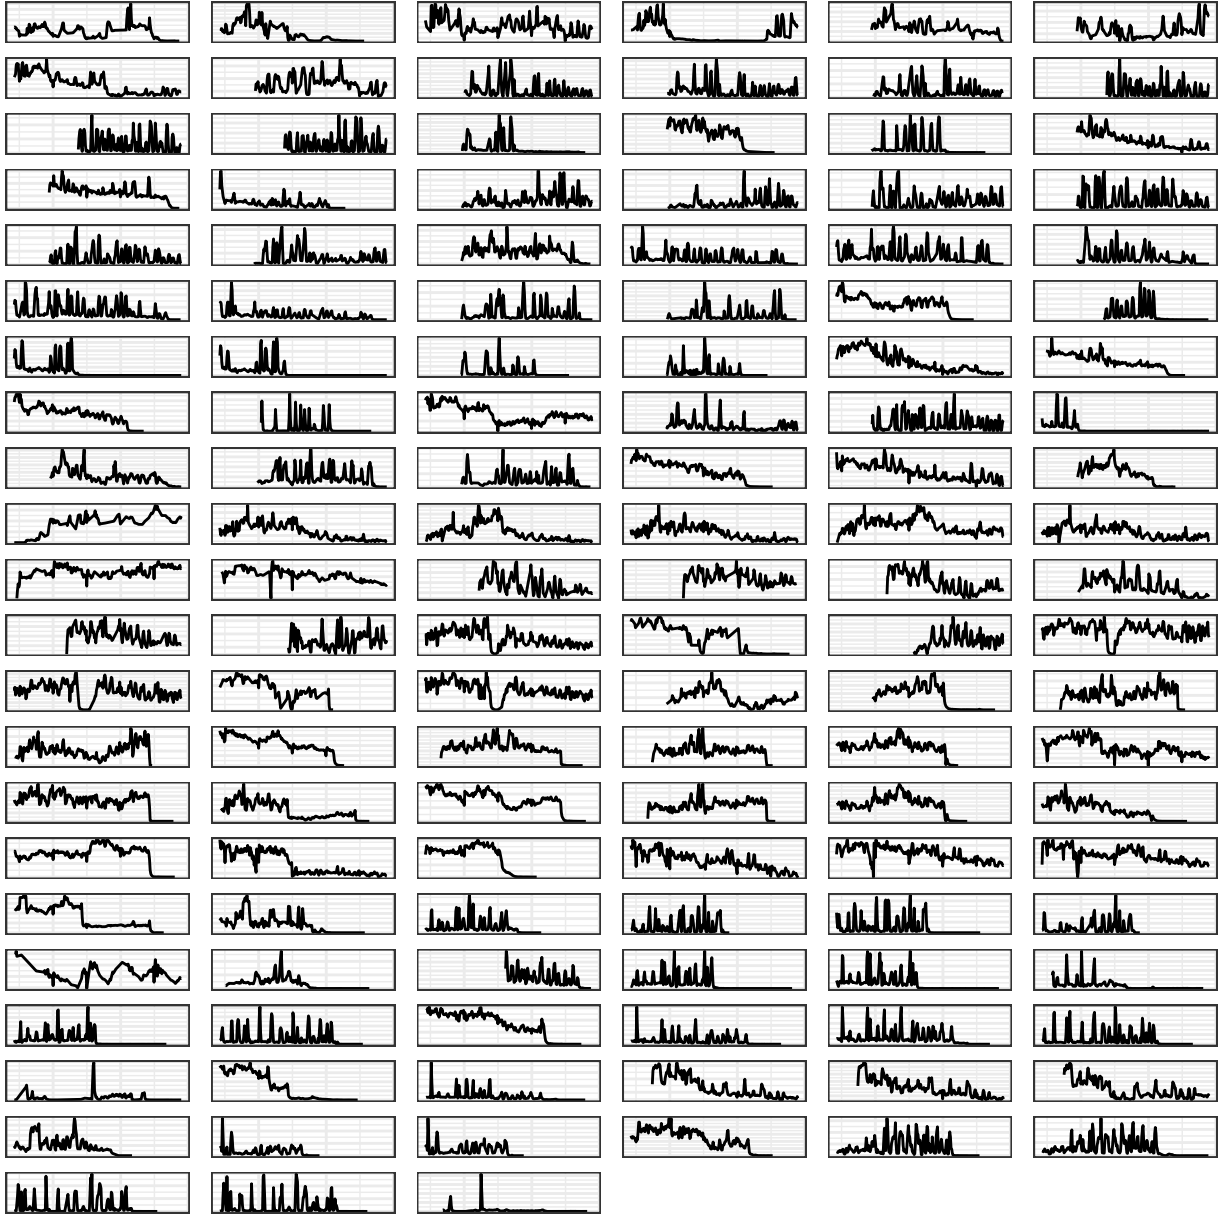


Figure 1: Data for a selection of SKUs.

## Forecasting

We'll limit forecasts for SKUs for which we have at least 52 measurements:

```
products %>% filter(count >= 52) -> products
```

Make a list of SKUs:

```
sort(unique(products$SKU)) -> SKUs
```

Select an SKU in list:

```
products %>% filter(SKU == SKUs[3]) -> product
```

There is missing data that is invisibly missing; the rows themselves are missing. Therefore we left-join to a column of dates covering the period such that the original data is padded with missing values:

```
begin <- range(product$Date)[1]
end <- range(product$Date)[2]
dates <- data.frame(Date = seq(from = begin, to = end, by = "week"))

left_join(dates, product, by = "Date") %>%
  mutate( # Expand-out data information
    year = as.numeric(format(Date, format = "%Y")),
    week = week(Date) # Week number
  ) -> product
```

Partition the data into training and test data:

```
require(caret)
```

```
## Loading required package: caret
```

```
## Warning: package 'caret' was built under R version 3.4.4
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
split.percent <- 0.85 # 15% holdout
```

```
p <- floor(split.percent * nrow(product))
```

```
q <- nrow(product) - p
```

```
print(c(p+q, p, q)) # Total number of data points and number in each partition
```

```
## [1] 205 174 31
```

```
in.train <- createTimeSlices(1:nrow(product), p, q)
```

```
df <- as.data.frame(product)
```

```
train <- df[unlist(in.train$train),]
```

```
test <- df[unlist(in.train$test),]
```

Start of reporting periods (both partitions):

```

start.train <- c(train$year[1], train$week[1])
start.test <- c(test$year[1], test$week[1])
print(start.train)

## [1] 2011    47

print(start.test)

## [1] 2015    12

Auto-fit an ARIMA model for an SKU, and interpolate missing values:

require(forecast)

## Loading required package: forecast
## Warning: package 'forecast' was built under R version 3.4.3

train.ts <- na.interp(
  ts(train$ValueSalesMLC, frequency = 52, start = start.train)
)
train.WeightedDistribution <- na.interp(
  ts(train$WeightedDistribution, frequency = 52, start = start.train)
)
train.PPSU <- na.interp(
  ts(train$PPSU, frequency = 52, start = start.train)
)
train.WDPriceCut <- na.interp(
  ts(train$WDPriceCut, frequency = 52, start = start.train)
)

xregs <- cbind(train.WeightedDistribution, train.PPSU, train.WDPriceCut)
# names.xregs <- colnames(xregs)
# xregs <- embed(xregs, 3)
#
# colnames(xregs) <- make.names(rep(names.xregs, 3), unique = TRUE)

arima.fit <- auto.arima(train.ts, trace = TRUE, xreg = xregs)

##
## Fitting models using approximations to speed things up...
##
## Regression with ARIMA(2,0,2)(1,1,1)[52] errors : -49.161
## Regression with ARIMA(0,0,0)(0,1,0)[52] errors : 336.3835
## Regression with ARIMA(1,0,0)(1,1,0)[52] errors : -27.34454
## Regression with ARIMA(0,0,1)(0,1,1)[52] errors : 273.3344
## ARIMA(0,0,0)(0,1,0)[52] : 421.7536
## Regression with ARIMA(2,0,2)(0,1,1)[52] errors : 212.8881
## Regression with ARIMA(2,0,2)(1,1,0)[52] errors : -50.66912
## Regression with ARIMA(1,0,2)(1,1,0)[52] errors : -50.75713
## Regression with ARIMA(1,0,1)(1,1,0)[52] errors : -53.05208
## Regression with ARIMA(0,0,0)(1,1,0)[52] errors : 269.7005
## ARIMA(1,0,1)(1,1,0)[52] : -54.88134
## ARIMA(1,0,1)(0,1,0)[52] : 220.7396
## ARIMA(1,0,1)(1,1,1)[52] : -52.59998
## ARIMA(0,0,1)(1,1,0)[52] : 173.7066
## ARIMA(2,0,1)(1,1,0)[52] : -50.55968

```

```
## ARIMA(1,0,0)(1,1,0) [52] : -29.59245
## ARIMA(1,0,2)(1,1,0) [52] : -52.69729
## ARIMA(0,0,0)(1,1,0) [52] : 276.1514
## ARIMA(2,0,2)(1,1,0) [52] : -52.37619
##
## Now re-fitting the best model(s) without approximations...
##
## ARIMA(1,0,1)(1,1,0) [52] : 274.3582
##
## Best model: Regression with ARIMA(1,0,1)(1,1,0) [52] errors
```

```
summary(arima.fit)
```

```
## Series: train.ts
## Regression with ARIMA(1,0,1)(1,1,0) [52] errors
##
## Coefficients:
##          ar1          ma1          sar1  train.WeightedDistribution  train.PPSU
##          0.9510  0.0110  -0.0015                        0.0218      0.1161
## s.e.      0.0296  0.2019   0.1390                        0.0177      0.6659
##          train.WDPriceCut
##                      0.0534
## s.e.                      0.0366
##
## sigma^2 estimated as 0.5062: log likelihood=-129.69
## AIC=273.38  AICc=274.36  BIC=293
##
## Training set error measures:
##                      ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.05890644 0.5809258 0.2824252 -2.245196 11.84714 0.1377238
##                      ACF1
## Training set -0.004902544
```

Forecast ahead for the next 12 weeks:

```
test.ts <- ts(test$ValueSalesMLC, frequency = 52, start = start.test)
```

```
test.WeightedDistribution <- na.interp(
  ts(test$WeightedDistribution, frequency = 52, start = start.test)
)
```

```
test.PPSU <- na.interp(
  ts(test$PPSU, frequency = 52, start = start.test)
)
```

```
test.WDPriceCut <- na.interp(
  ts(test$WDPriceCut, frequency = 52, start = start.test)
)
```

```
xregs <- cbind(test.WeightedDistribution, test.PPSU, test.WDPriceCut)
```

```
# xregs <- cbind(test.WeightedDistribution, test.PPSU, test.WDPriceCut)
# names.xregs <- colnames(xregs)
# xregs <- embed(xregs, 3)
#
# colnames(xregs) <- make.names(rep(names.xregs, 3), unique = TRUE)
```



```
arima.forecast <- forecast(arima.fit, h = q, xreg = xregs)
```

Plot the forecast:

```
train.df <- as.data.frame(time(train.ts))
names(train.df) <- "x"
train.df$y <- as.vector(train.ts)
train.df$Partition <- "Train"


test.df <- as.data.frame(time(test.ts))
names(test.df) <- "x"
test.df$y <- as.vector(test.ts)
test.df$Partition <- "Test"

points.dat <- suppressWarnings(bind_rows(test.df, train.df))

autoplot(arima.forecast) +
  ylab("ValueSalesMLC") +
  geom_point(data = points.dat, aes(x = x, y = y, colour = Partition)) +
  scale_color_manual(values = c("Train" = 'Black', 'Test' = 'Red'))
```

```
## Warning: Removed 17 rows containing missing values (geom_point).
```

### Forecasts from Regression with ARIMA(1,0,1)(1,1,0)[52] errors



```
accuracy(arima.forecast)
```

```
##
## Training set  ME      RMSE      MAE      MPE      MAPE      MASE
```

```
##                               ACF1
## Training set -0.004902544
```

We perform a statistical significant test of the coefficients of the fitted model; if the  $p$ -value is less than 5% we take this as evidence that the coefficient is statistically significant. The size of the coefficient provides a measure of the corresponding variable's importance in the model. **This addresses the challenge of determining the performance drivers for a specific product:**

```
require(lmtest)
```

```
## Loading required package: lmtest
## Warning: package 'lmtest' was built under R version 3.4.3
## Loading required package: zoo
## Warning: package 'zoo' was built under R version 3.4.4
##
## Attaching package: 'zoo'
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric
```

```
coeftest(arima.fit)
```

```
##
## z test of coefficients:
##
##              Estimate Std. Error z value Pr(>|z|)
## ar1              0.9509782  0.0295864  32.1424  <2e-16 ***
## ma1              0.0109740  0.2018835   0.0544   0.9567
## sar1             -0.0015001  0.1390271  -0.0108   0.9914
## train.WeightedDistribution 0.0218041  0.0176892   1.2326   0.2177
## train.PPSU        0.1160875  0.6658717   0.1743   0.8616
## train.WDPriceCut   0.0534070  0.0366053   1.4590   0.1446
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
print(arima.fit$aic) # Print Akaike information criterion (AIC)
```

```
## [1] 273.3757
```

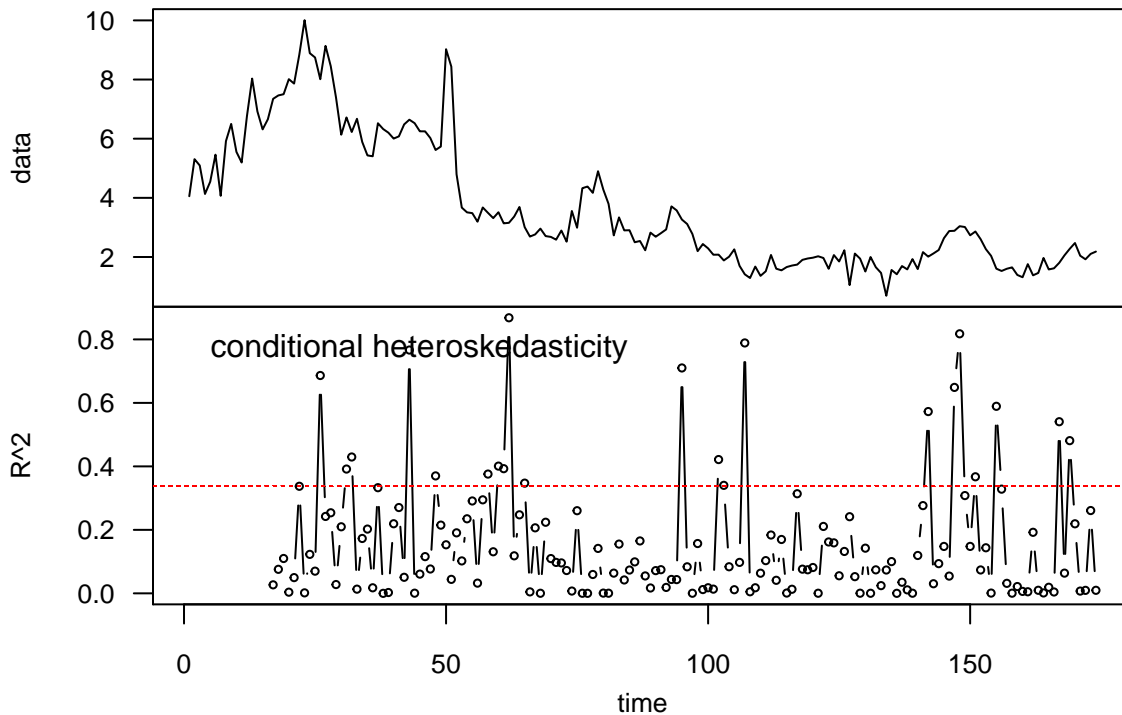
```
require(earlywarnings)
```

```
## Loading required package: earlywarnings
## Warning: package 'earlywarnings' was built under R version 3.4.3
## Loading required package: moments
## Warning: package 'moments' was built under R version 3.4.1
## Loading required package: tgp
## Warning: package 'tgp' was built under R version 3.4.3
## Loading required package: tseries
## Warning: package 'tseries' was built under R version 3.4.4
##
## earlywarnings Copyright (C) 2011-2013 Vasilis Dakos and Leo Lahti
```

```
## This program comes with ABSOLUTELY NO WARRANTY.
## This is free software, and you are welcome to redistribute it under the FreeBSD open source license.
```

```
out <- quiet(bdstest_ews(as.data.frame(train.ts), ARMAoptim = FALSE))
```

```
out <- quiet(ch_ews(as.data.frame(train.ts)))
```



```
out <- quiet(ddjnonparam_ews(as.data.frame(train.ts)))
```

```
out <- quiet(generic_ews(as.data.frame(train.ts)))
```

```
res <- livpotential_ews(as.data.frame(train.ts))
```

```
out <- quiet(qda_ews(as.data.frame(train.ts)))
```

```
#plot(out$potential.plot$plot)
```

```
out <- quiet(sensitivity_ews(as.data.frame(train.ts)))
```

```
out <- quiet(surrogates_ews(as.data.frame(train.ts)))
```

Wrap up everything we've just done for one SKU in a *crystal ball* function that we can run for any SKU:

```
crystal.ball <- function(SKU.num, split.percent, verbose = FALSE) {
  products %>% filter(SKU == SKUs[SKU.num]) -> product

  begin <- range(product$Date)[1]
  end <- range(product$Date)[2]
  dates <- data.frame(Date = seq(from = begin, to = end, by = "week"))
```

```

left_join(dates, product, by = "Date") %>%
  mutate( # Expand-out data information
    year = as.numeric(format(Date, format = "%Y")),
    week = week(Date) # Week number
  ) -> product

p <- floor(split.percent * nrow(product))
q <- nrow(product) - p

in.train <- createTimeSlices(1:nrow(product), p, q)

df <- as.data.frame(product)
train <- df[unlist(in.train$train),]
test <- df[unlist(in.train$test),]

start.train <- c(train$year[1], train$week[1])
start.test <- c(test$year[1], test$week[1])

train.ts <- na.interp(
  ts(train$ValueSalesMLC, frequency = 52, start = start.train)
)
train.WeightedDistribution <- na.interp(
  ts(train$WeightedDistribution, frequency = 52, start = start.train)
)
train.PPSU <- na.interp(
  ts(train$PPSU, frequency = 52, start = start.train)
)
train.WDPriceCut <- na.interp(
  ts(train$WDPriceCut, frequency = 52, start = start.train)
)
train.xregs <- cbind(train.WeightedDistribution, train.PPSU, train.WDPriceCut)

test.ts <- ts(test$ValueSalesMLC, frequency = 52, start = start.test)

test.WeightedDistribution <- na.interp(
  ts(test$WeightedDistribution, frequency = 52, start = start.test)
)
test.PPSU <- na.interp(
  ts(test$PPSU, frequency = 52, start = start.test)
)
test.WDPriceCut <- na.interp(
  ts(test$WDPriceCut, frequency = 52, start = start.test)
)
test.xregs <- cbind(test.WeightedDistribution, test.PPSU, test.WDPriceCut)

arima.fit <- auto.arima(train.ts, trace = FALSE, xreg = train.xregs)

# res = AnomalyDetectionVec(as.vector(train.ts), max_anoms=0.02, period = 52, direction='both', plot=
# plot(res$plot)

if(verbose) {
  print(coefest(arima.fit))
  print(arima.fit$aic) # Print Akaike information criterion (AIC)

```

```

}

arima.forecast <- forecast(arima.fit, h = q, xreg = test.xregs)

train.df <- as.data.frame(time(train.ts))
names(train.df) <- "x"
train.df$y <- as.vector(train.ts)
train.df$Partition <- "Train"

test.df <- as.data.frame(time(test.ts))
names(test.df) <- "x"
test.df$y <- as.vector(test.ts)
test.df$Partition <- "Test"

points.dat <- suppressWarnings(bind_rows(test.df, train.df))


autoplot(arima.forecast) +
  ylab("ValueSalesMLC") +
  geom_point(data = points.dat, aes(x = x, y = y, colour = Partition)) +
  scale_color_manual(values = c("Train" = 'Black', 'Test' = 'Red'))
}

crystal.ball(5, 0.85, verbose = TRUE)

##
## z test of coefficients:
##
##               Estimate Std. Error  z value  Pr(>|z|)
## ar1              0.8137498  0.0746154  10.9059 < 2.2e-16 ***
## ma1             -1.2867231  0.1119447 -11.4943 < 2.2e-16 ***
## ma2              0.3049328  0.1043563   2.9220  0.003478 **
## train.WeightedDistribution  0.0720436  0.0041545  17.3412 < 2.2e-16 ***
## train.PPSU        0.0145742  0.0113041   1.2893  0.197302
## train.WDPriceCut   0.0148697  0.0033346   4.4592 8.227e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## [1] 24.98226
## Warning: Removed 19 rows containing missing values (geom_point).

```

## Forecasts from Regression with ARIMA(1,1,2) errors




Let's test our crystal ball function on some randomly-selected SKUs:

```
random.SKUs <- sample(length(SKUs), size = 6, replace = FALSE)
lapply(random.SKUs, function(N) crystal.ball(N, 0.85))
```


```
## [[1]]
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```

Forecasts from Regression with ARIMA(1,1,1) errors




```
##  
## [[2]]
```




```
##
## [[3]]
## Warning: Removed 29 rows containing missing values (geom_point).
```






```
##
## [[4]]
## Warning: Removed 17 rows containing missing values (geom_point).
```

Forecasts from Regression with ARIMA(1,0,1)(1,1,0)[52] errors




##  
## [[5]]

Forecasts from Regression with ARIMA(0,0,2) errors



```
##
## [[6]]
## Warning: Removed 16 rows containing missing values (geom_point).
```

## Forecasts from Regression with ARIMA(0,1,0)(1,1,0)[52] errors



## What could have been done better

We could introduce lagged versions of the explanatory variables into the model, and we could select the best model using AIC – we didn't do this here, but it would be reasonably trivial to implement with more time.

```
sessionInfo()
```

```
## R version 3.4.0 (2017-04-21)
## Platform: x86_64-w64-mingw32/x64 (64-bit)
## Running under: Windows 10 x64 (build 14393)
##
## Matrix products: default
##
## locale:
## [1] LC_COLLATE=English_United Kingdom.1252
## [2] LC_CTYPE=English_United Kingdom.1252
## [3] LC_MONETARY=English_United Kingdom.1252
## [4] LC_NUMERIC=C
## [5] LC_TIME=English_United Kingdom.1252
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] earlywarnings_1.0.59 tseries_0.10-43      tgp_2.4-14
## [4] moments_0.14         lmtest_0.9-35        zoo_1.8-1
```

```

## [7] forecast_8.2          caret_6.0-78          lattice_0.20-35
## [10] bindrcpp_0.2          digest_0.6.15         lubridate_1.7.3
## [13] readxl_1.0.0          forcats_0.3.0         stringr_1.3.0
## [16] dplyr_0.7.4           purrr_0.2.4           readr_1.1.1
## [19] tidyr_0.8.0           tibble_1.4.2          ggplot2_2.2.1
## [22] tidyverse_1.2.1
##
## loaded via a namespace (and not attached):
## [1] nlme_3.1-131.1        xts_0.10-2            dimRed_0.1.0
## [4] httr_1.3.1           rprojroot_1.3-2       tools_3.4.0
## [7] backports_1.1.2       R6_2.2.2              KernSmooth_2.23-15
## [10] rpart_4.1-11          nortest_1.0-4         lazyeval_0.2.1
## [13] colorspace_1.3-2     nnet_7.3-12           withr_2.1.2
## [16] tidyselect_0.2.4     mnormt_1.5-5          curl_3.1
## [19] compiler_3.4.0       cli_1.0.0             rvest_0.3.2
## [22] xml2_1.2.0           labeling_0.3           scales_0.5.0
## [25] sfsmisc_1.1-2        DEoptimR_1.0-8        fracdiff_1.4-2
## [28] psych_1.7.8          robustbase_0.92-8     quadprog_1.5-5
## [31] foreign_0.8-69       rmarkdown_1.9         pkgconfig_2.0.1
## [34] htmltools_0.3.6      maps_3.2.0            highr_0.6
## [37] TTR_0.23-3           rlang_0.2.0           ddtalpha_1.3.1.1
## [40] quantmod_0.4-12      rstudioapi_0.7        bindr_0.1.1
## [43] jsonlite_1.5         ModelMetrics_1.1.0   magrittr_1.5
## [46] dotCall64_0.9-5.2   Matrix_1.2-12         mape_1.4-7
## [49] Rcpp_0.12.16         munsell_0.4.3         stringi_1.1.7
## [52] yaml_2.1.18          MASS_7.3-49           plyr_1.8.4
## [55] recipes_0.1.2        grid_3.4.0            parallel_3.4.0
## [58] crayon_1.3.4         haven_1.1.1           splines_3.4.0
## [61] hms_0.4.2           knitr_1.20            pillar_1.2.1
## [64] boot_1.3-20          reshape2_1.4.3        codetools_0.2-15
## [67] stats4_3.4.0         CVST_0.2-1            glue_1.2.0
## [70] evaluate_0.10.1     modelr_0.1.1          spam_2.1-2
## [73] foreach_1.4.4        cellranger_1.1.0      gtable_0.2.0
## [76] kernlab_0.9-25       assertthat_0.2.0     DRR_0.0.3
## [79] gower_0.1.2         prodlim_1.6.1         broom_0.4.3
## [82] Kendall_2.2          class_7.3-14          survival_2.41-3
## [85] timeDate_3043.102   RcppRoll_0.2.2        som_0.3-5.1
## [88] iterators_1.0.9     fields_9.6            cluster_2.0.6
## [91] lava_1.6             ipred_0.9-6

```